

Info Fusion - Sage Lab Synthesis: Improving antibiotic resistance prediction with mutli-modal input

Jackson Michaels



Project Outline

Background - Bio

Genetic Vocab:

- DNA strand of **base pairs** aka **nucleotide pairs**
- sets of three (**codons**) are encoded as **amino acids**
- amino acids come together to construct **proteins**
- generally, **one protein per gene** for TB
- Regions can be **genic** (inside a gene) or **intergenic** (outside a gene)
 - the latter are of particular interest because they regulate and promote gene expression

Project Vocab:

- previous works have identified sets of relevant regions in the genome referred to as **Loci**
- **Scope**, A locus can be entirely genic, entirely intergenic, or span across both.

Background - Project

Goal: take existing work from SAGE lab on antibiotic resistance prediction in TB and improve on performance and interpretability by incorporating multimodal insights from Info Fusion lab

Task Details:

- Original input data is presented as set of loci, 5 X ~7,500 vector of base pairs for single drug
 - these loci can be genic regions or intergenic
 - typically one hot encoded (*A, C, T, G, gap*)
 - covers 2-3 genes and intergenic regions
- task is to predict resistance to eleven antibiotic drugs
 - multi-drug input is concatenated single-drug inputs

```
GTCGAGCACGGTGCAGGCCAATGCCGCGTCCATAAAGACGTCCATCACACTCATGCGTCGACGGTAGC
CAGCTCGTGTGCCACGTCCGGTTACGGCGCAGGTATTTCTGCAAGGTAGTGTGAGTACGCAGCTGCCATCG
```

Modalities - DNA

Top Choices:

- Regulatory regions
- more than 2-3 proteins

DNA-Derived:

- one-hot array / ref-alt encoding
- foundation model embedding
- K-mer frequency (31-mers used in literature)

```
>NC_000962.3 NC_000962.3:314309..314854 (-  
strand) class=CDS length=546
```

```
GTGCACACCCAGGTACACACGGCCCGCCTGGTCCAC  
ACCGCCGATCTTGACAGCGAGACCCGCCAGGACATC  
CGTCAGATGGTCACCGGCGCGTTTGCCGGTGACTTCA  
CCGAGACCGACTGGGAGCACACGCTGGGTGGGATGC  
ACGCCCTGATCTGGCATCACGGGGCGATCATCGCGC  
ATGCCGCGGTGATCCAGCGGCGACTGATCTACCGCG  
GCAACGCGCTGCGCTGCGGGTACGTCTGAAGGCGTTG  
CGGTGCGGGCGGACTGGCGGGGCCAACGCCTGGTG  
AGCGCGCTGTTGGACGCCGTCGAGCAGGTGATGCGC  
GGCGCTTACCAGCTCGGAGCGCTCAGTTCCTCGGCG  
CGGGCCCGCAGACTGTACGCCTCACGCGGCTGGCTG  
CCCTGGCACGGCCCGACATCGGTACTGGCACCAACC  
GGTCCAGTCCGTACACCCGATGACGACGGAACGGTG  
TTCGTCTGCCCATCGACATCAGCCTGGACACCTCGG  
CGGAGCTGATGTGCGATTGGCGCGCGGGGCGACGTCT  
GGTAA
```

Example Loci: Rv0262c

Modalities - Regulatory

Regulatory Regions:

- high-value, mechanistically-motivated modality that current CNN baselines under-use
- handled as separate modality from DNA

Description:

- regulates the downstream gene(s) it precedes
 - e.g. promoters / operators
- regulates expression or lack of expression for previous genes

Motivation:

- Resistance can arise without changing the protein at all, by a promoter mutation that over-expresses the gene

Modalities - Protein

Description:

- Translation from DNA sequence
- effectively 3-mers

Representations:

- one-hot array / ref-alt encoding of amino-acids
- foundation model embedding

Notes:

- another top choice is to simply include more than 2-3 in the inference

```
>Mycobacterium tuberculosis
```

```
H37Rv|Rv0262c|aac
```

```
VHTQVHTARLVHTADLDSETRQDIRQMVTGAFAGDFT  
ETDWEHTLGGMHALIWHHGAIIAHA AVIQRRLIYRGNA  
LRCGYVEGVAVRADWRGQRLVSALLDAVEQVMRGAYQ  
LGALSSSARARRLYASRGWLPWHGPTSVLAPTGPVRTP  
DDDGTVFVLPIDISLDTSAELMCDWRAGDVW
```

Example Loci: Rv0262c, amino-acid sequence

Modalities - 3D Structure

Description:

- spatially describe distance of base-pairs in folded protein

Representations:

- distance matrix (e.g. alpha fold)
- Density vector (mine)
- Shrake-Rupley algorithm

Notes:

- only useful if two proteins interact



Modalities - Molecular Characteristics

Description:

- Chemical properties of molecules in protein separate from function
- Represents how individual molecules behave and interact
- Used in (*Kulkarni et al*)

Features:


- Molecular weight
- Isoelectric point
- Hydro / Lipophilicity
- Refractivity

Notes:

- derived from RDKit Library



Open-Source Cheminformatics
and Machine Learning



```
getMolDescriptors(doravirine)
```

```
{'MaxEStateIndex': 13.412553309006833,  
'MinEStateIndex': -4.871620672188628,  
'MaxAbsEStateIndex': 13.412553309006833,  
'MinAbsEStateIndex': 0.045220418860841605,  
'qed': 0.6914051268589834,  
'MolWt': 425.754,  
...  
'fr_thiophene': 0,  
'fr_unbrch_alkane': 0,  
'fr_urea': 0}
```

Modalities - Lineage

Description:


- evolutionary lineage of specific TB strain

Representation:

- 1D binary vector, N X N matrix, tree structure

Notes:

- Literature is split on using lineage as its commonly a confounder
- I will ablate with and without lineage




Model Design

Models - Literature

Trends:

- CNN models consistently outperform other architectures
 - particularly on DNA sequence tasks
- DNA-sequence models consistently outperform protein-based models, likely because they capture crucial non-coding regulatory regions.
- Despite their complexity, language models struggle with second-line drugs and sequence-position averaging, largely due to data bottlenecks.




Models - Cross Attention

Description:

- to retain 1D CNN performance the model uses cross attention treating DNA as the query and other modalities as the key / value
- Should allow CNN insights to guide inference only retrieving information as needed

Cross-Attention Fusion CNN (asymmetric: DNA = Query, biophysical = Key/Value)

Representative scenario — Rifampicin (loci rpoB + rpoC): DNA one-hot 5x240, biophysical 3 channels, protein lengths K = 78 / 66




Models - Early Fusion

Description:

- concatenate all modalities into single vector
- pass concat vector into model as normal
- Fewer parameters and simpler implementation but lacking multi-modal nuance

Early-Fusion CNN (channel-stack ablation)

Representative scenario — Rifampicin (loci rpoB + rpoC): DNA one-hot 5x240, biophysical 3 channels, protein lengths K = 78 / 66




Models - Late Fusion

Description:

- separate stacks for each modality
- allows different models to process different modalities
 - e.g. CNN for DNA and transformer for 3D distance
- requires more parameters but more robust
- can be trained with modality drop-out

Late-Fusion CNN (default / Kulkarni et al. 2026)

Representative scenario — Rifampicin (loci rpoB + rpoC): DNA one-hot 5x240, biophysical 3 channels, protein lengths K = 78 / 66




Models - Late Fusion

Description:

- separate stacks for each modality
- allows different models to process different modalities
 - e.g. CNN for DNA and transformer for 3D distance
- requires more parameters but more robust
- can be trained with modality drop-out

Late-Fusion CNN (default / Kulkarni et al. 2026)

Representative scenario — Rifampicin (loci rpoB + rpoC): DNA one-hot 5x240, biophysical 3 channels, protein lengths K = 78 / 66



Interpretability

Interp - SHAP

Description:

- Open Source
- mechanistic interpretability library for input attribution
- commonly used in Comp Bio papers

